



上海科技大学
ShanghaiTech University

CS283: Robotics Spring 2025: SLAM I


Sören Schwertfeger / 师泽仁

ShanghaiTech University

ICP

ICP: Iterative Closest Points Algorithm


- Align two partially-overlapping point sets (2D or 3D)
- Given initial guess for relative transform
- Warning: Using 3D ICP for 2D data may mirror the data (e.g. 180 degree roll)!
 - Use 2D ICP!
- ROS: Point Cloud Library (PCL)



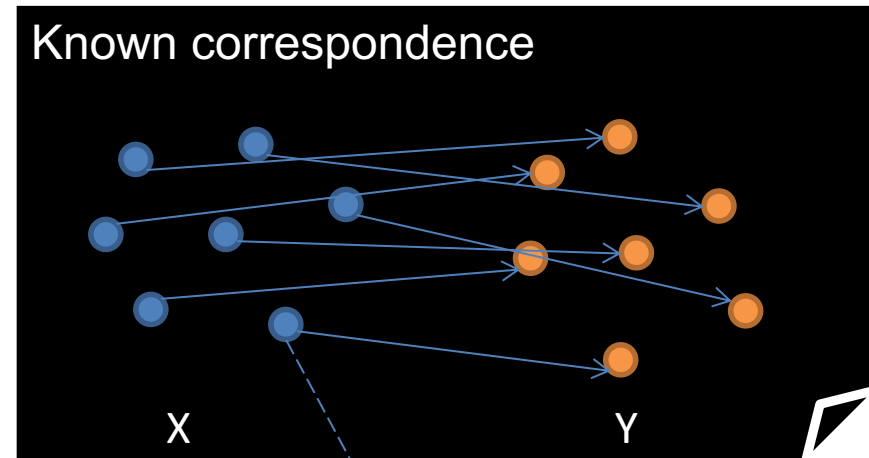

Horn's method

Material by Toru Tamaki, Miho Abe,
Bisser Raytchev, Kazufumi Kaneda


- Input
 - Two point sets: X and Y
- Output
 - Rotation matrix R
 - Translation vector t
 - Fitting error



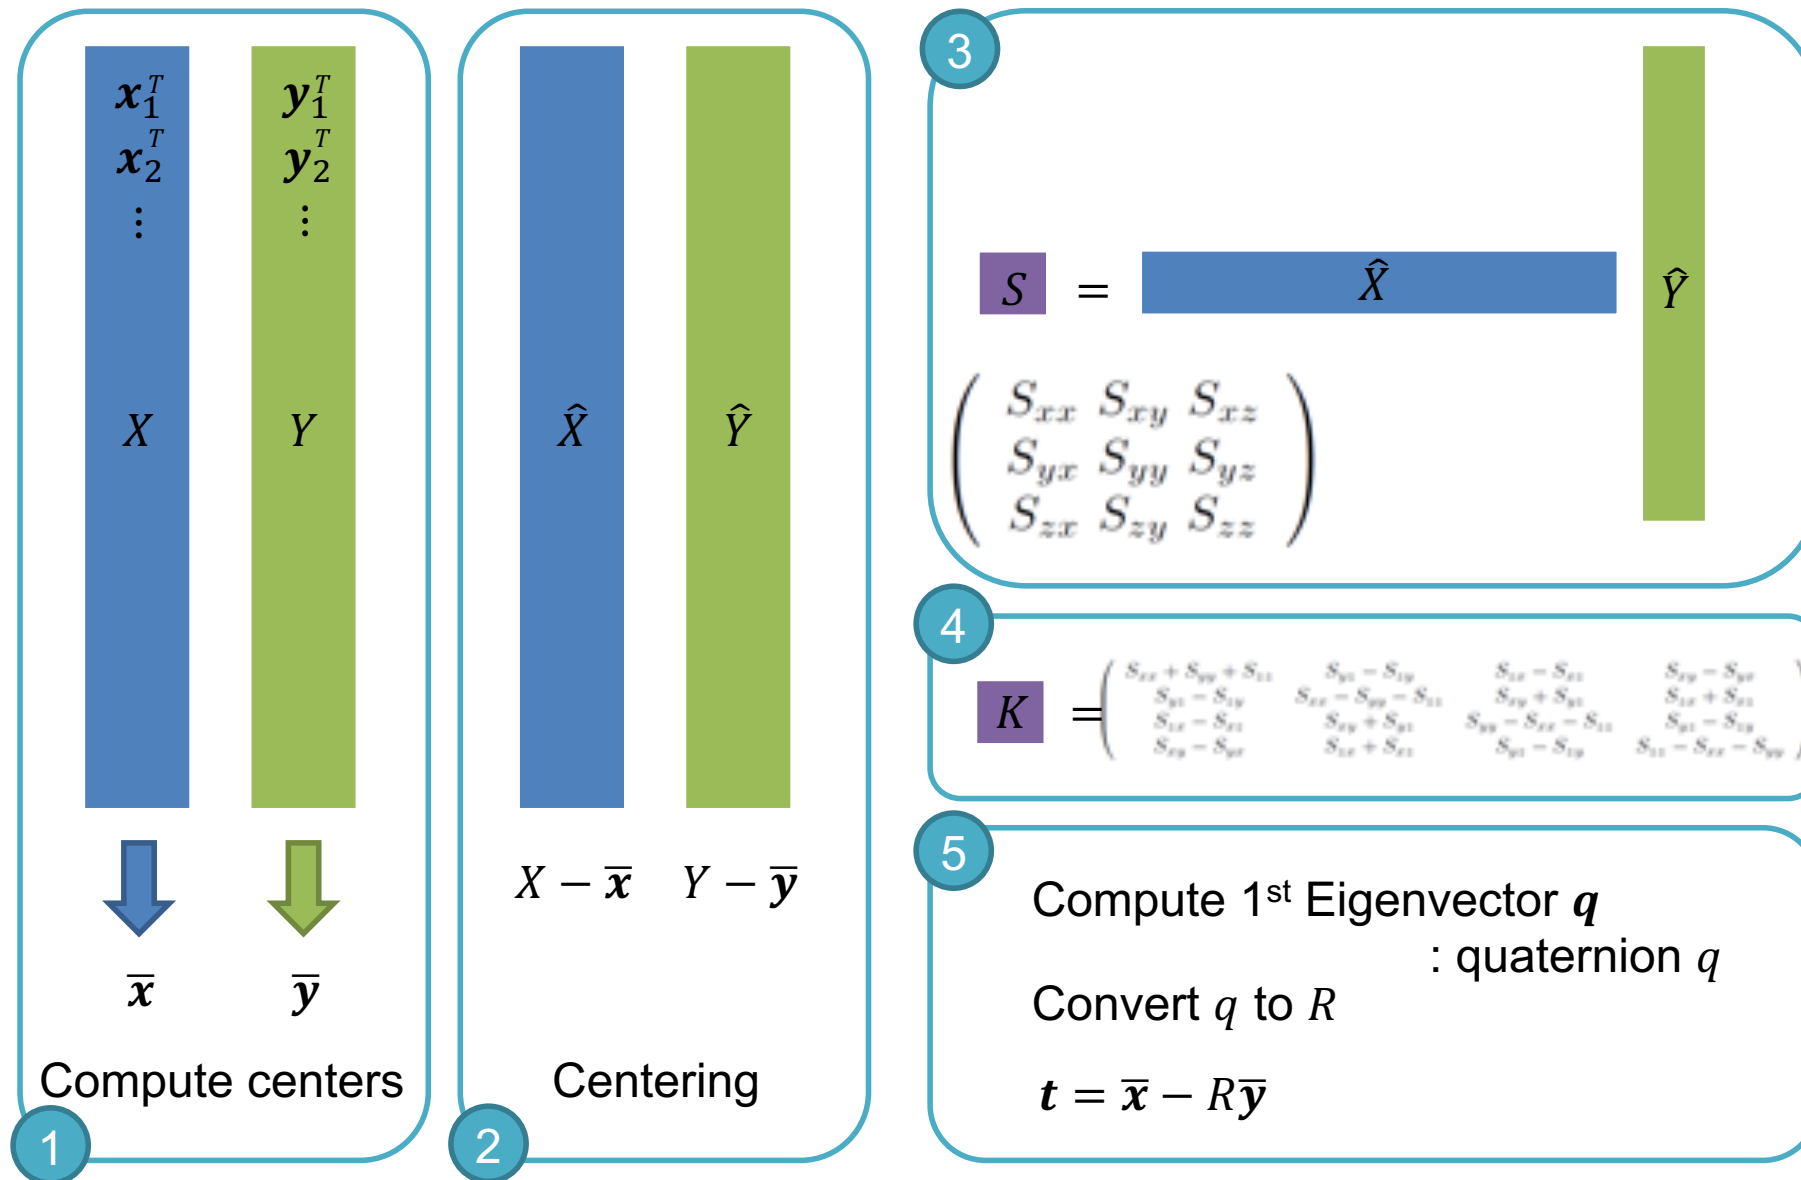
Horn's method: correspondence is known.



$$x_1 = (x_{1x}, x_{1y}, x_{1z})^T$$




Horn's method: correspondence is known.




Aligning 3D Data

- How to find correspondences: User input? Feature detection? Signatures?
- Alternative: assume **closest** points correspond




Aligning 3D Data


- Converges if starting position “close enough”



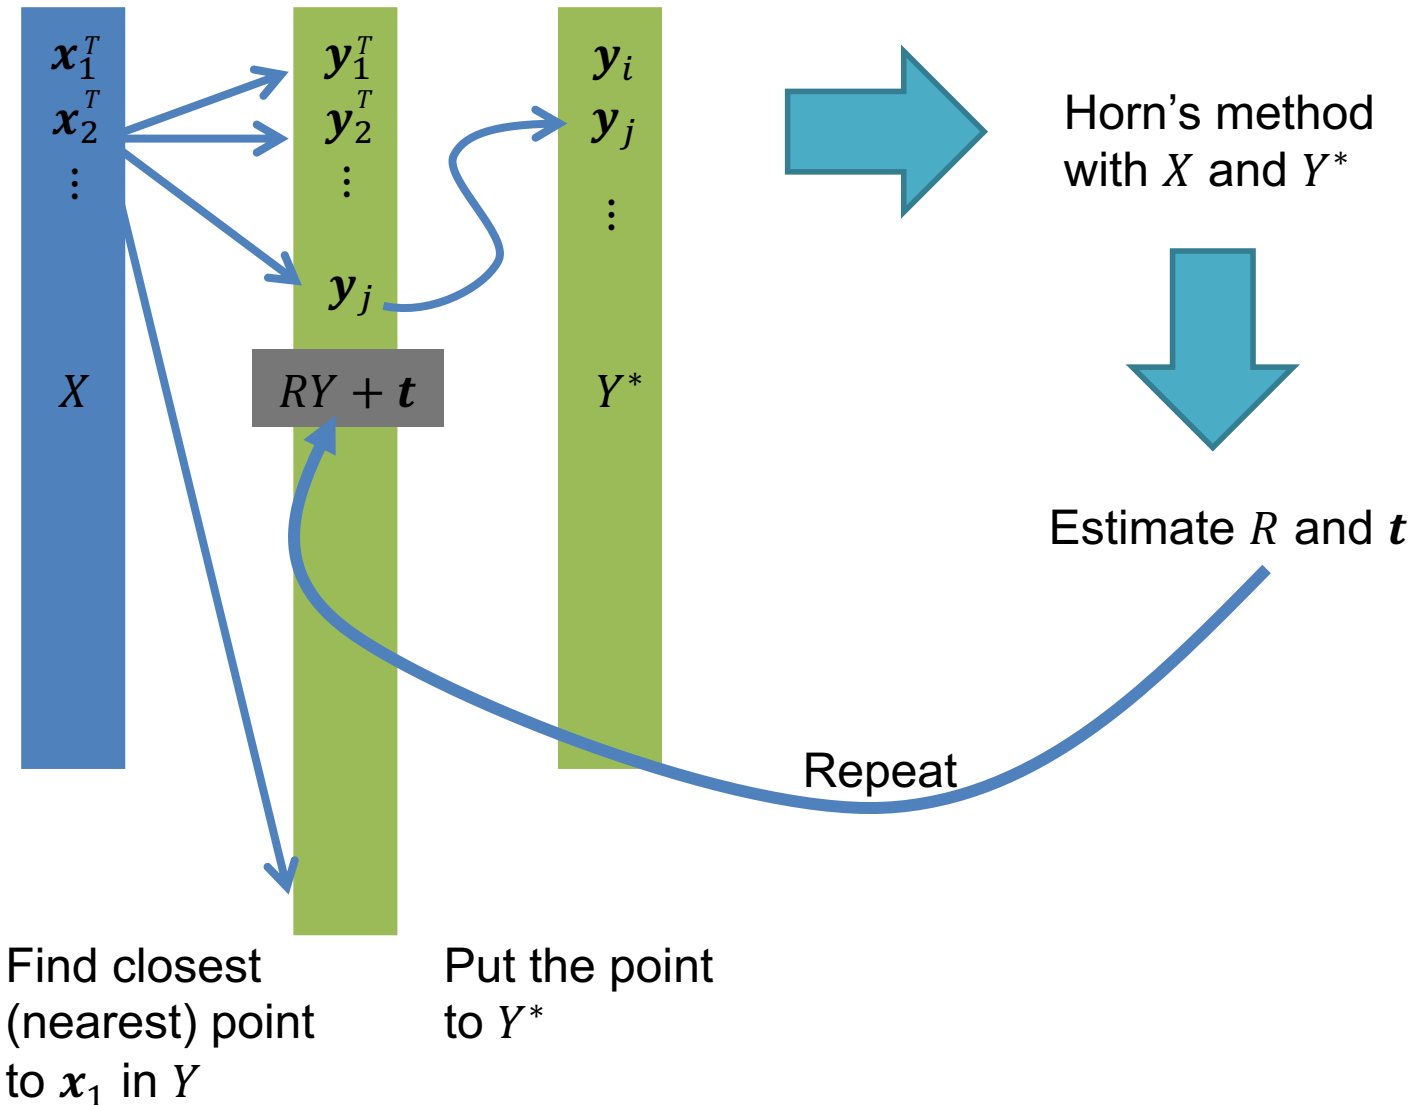
ICP: correspondence is unknown.



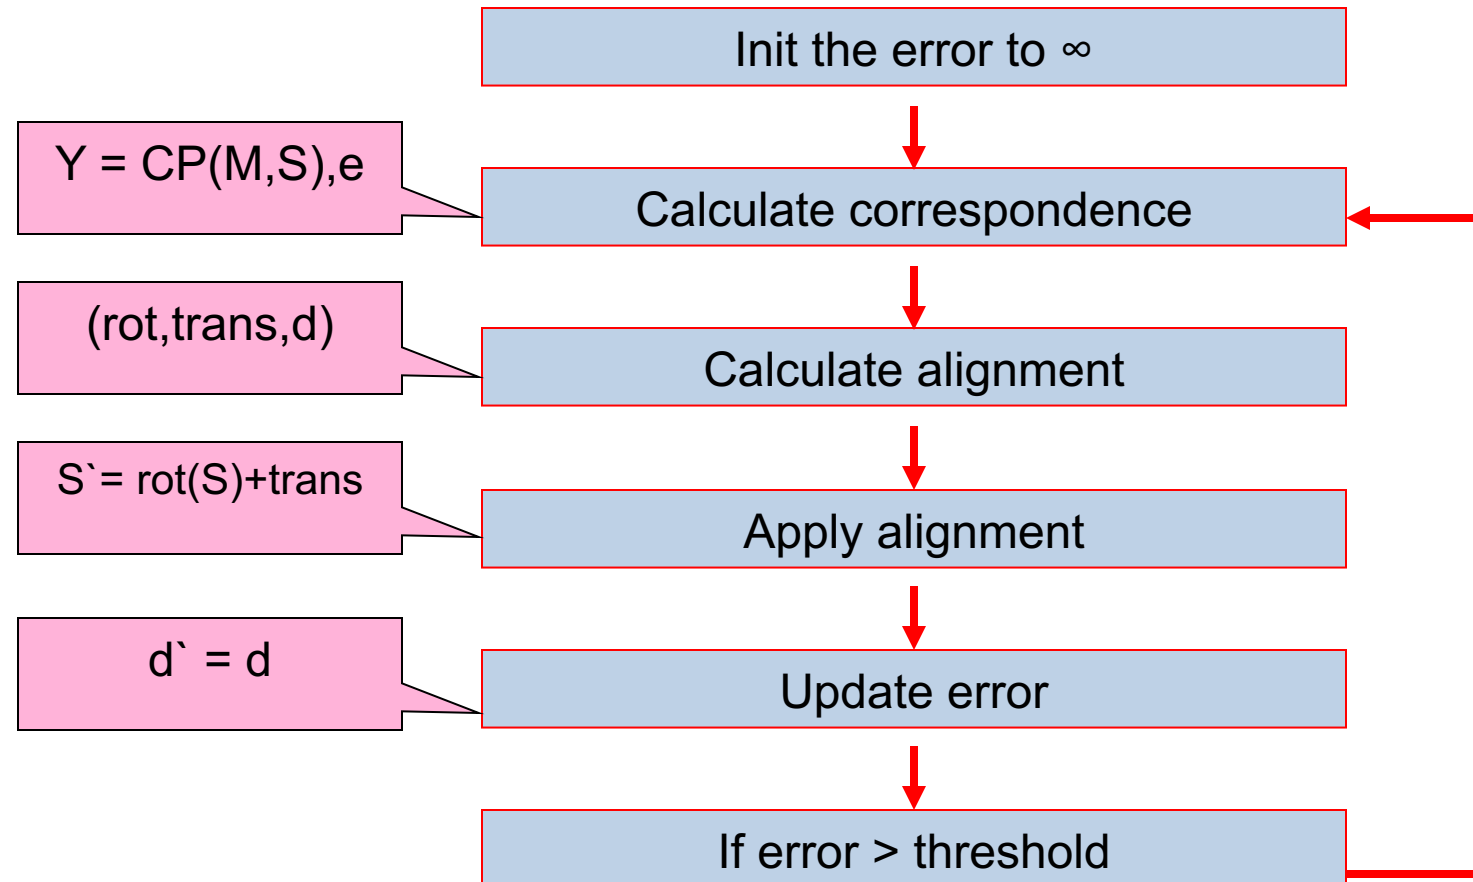
ICP: correspondence is unknown.



ICP: correspondence is unknown.



The Algorithm



The Algorithm

```
function ICP(Scene,Model)
begin
 $E' \leftarrow +\infty$ ;
(Rot,Trans)  $\leftarrow$  Initialize-Alignment(Scene,Model);
repeat
     $E \leftarrow E'$ ;
    Aligned-Scene  $\leftarrow$  Apply-Alignment(Scene,Rot,Trans);
    Pairs  $\leftarrow$  Return-Closest-Pairs(Aligned-Scene,Model);
    (Rot,Trans, $E'$ )  $\leftarrow$  Update-Alignment(Scene,Model,Pairs,Rot,Trans);
Until  $|E' - E| < \text{Threshold}$ 
return (Rot,Trans);
end
```

Convergence Theorem

- The ICP algorithm always converges monotonically to a local minimum with respect to the MSE distance objective function.

Time analysis

Each iteration includes 3 main steps

A. Finding the closest points :

$O(N_M)$ per each point

$O(N_M * N_S)$ total.

B. Calculating the alignment: $O(N_S)$

C. Updating the scene: $O(N_S)$

Optimizing the Algorithm

The best match/nearest neighbor problem :

Given **N** records each described by **K** real values (attributes), and a dissimilarity measure **D**, find the **m** records closest to a query record.

Optimizing the Algorithm

- K-D Tree :

Construction time: $O(kn \log n)$

Space: $O(n)$

Region Query : $O(n^{1-1/k} + k)$

Time analysis

Each iteration includes 3 main steps

A. Finding the closest points :

$O(N_M)$ per each point

$O(N_M \log N_S)$ total.

B. Calculating the alignment: $O(N_S)$

C. Updating the scene: $O(N_S)$

Further optimization: Approximate k-d tree search

ICP Variants

- Variants on the following stages of ICP have been proposed:
 1. **Selecting** sample points (from one or both point clouds)
 2. **Matching** to points to a plane or mesh
 3. **Weighting** the correspondences
 4. **Rejecting** certain (outlier) point pairs
 5. Assigning an **error metric** to the current transform
 6. **Minimizing** the error metric w.r.t. transformation
- Can analyze various aspects of performance:
 - Speed
 - Stability
 - Tolerance to noise and/or outliers
 - Maximum initial misalignment

Rejecting Pairs


- Corresponding points with point to point distance higher than a given threshold.
- Rejection of worst $n\%$ pairs based on some metric.
- Pairs containing points on end vertices.
- Rejection of pairs whose point to point distance is higher than $n^*\sigma$.
- Rejection of pairs that are not consistent with their neighboring pairs [Dorai 98] :

(p_1, q_1) , (p_2, q_2) are inconsistent iff

$$|Dist(p_1, p_2) - Dist(q_1, q_2)| > threshold$$


Rejecting Pairs

Distance thresholding




Rejecting Pairs

Points on end vertices




Rejecting Pairs

Inconsistent Pairs



BLAM: ICP in action






DEFINITION OF SLAM

What is SLAM?

- **Localization**: inferring location given a map
- **Mapping**: inferring a map given locations
- **SLAM**: learning a map and locating the robot simultaneously
- SLAM has long been regarded as a chicken-and-egg problem:
 - a map is needed for localization and
 - a pose estimate is needed for mapping




Material derived from Wolfram Burgard:

<http://ais.informatik.uni-freiburg.de/teaching/ss20/robotics/slides/13-slam.pdf>


SLAM Applications

- At home: vacuum cleaner, lawn mower
- Air: surveillance with unmanned air vehicles
- Underwater: reef monitoring
- Underground: exploration of mines
- Space: terrain mapping for localization
-




Why is SLAM a Hard Problem?

- Robot path and map are both unknown
- Errors in map and pose estimates correlated



Why is SLAM a Hard Problem?

- The mapping between observations and landmarks is unknown
- Picking wrong data associations can have catastrophic consequences (divergence)



Overview of SLAM Methods

- Camera

- Feature-Based Methods
 - MonoSLAM
 - PTAM
 - ORB-SLAM
- Direct Methods
 - DTAM
 - LSD-SLAM
 - DSO
- Semi-Direct Methods
 - SVO
- Others
 - PoseNet
 - CNN-SLAM
 - ...

- Laser

- Pose Graph
 - Cartographer
 - Karto-SLAM
 - Hector-SLAM
 - BLAM
 - LIO
- Particle Filter
 - FastSLAM
 - Gmapping
- Extended Kalman Filter
 - EKF-SLAM
 - LINS
- Others
 - LOAM
 - IMLS-SLAM
 - ...

SLAM Front-end & Back-end

- Front-end
 - calculate relative poses between several frames/ to map
 - scan matching
 - image registration
 - ...
 - estimate absolute poses
 - construct the local map
- Back-end
 - optimize the absolute poses and mapping
 - only if a loop was closed



FRONT END – LASER - ICP

FRONT END - CAMERA

Methods

- **Feature-based Methods**

- SIFT
- ORB (ORB-SLAM)
- BRISK
- AKAZE



- **Direct Methods**

- Optical Flow
- Inverse Depth (LSD-SLAM)
- Fourier-Mellin Transform

- **Semi-Direct Methods**

- SVO

more details in the next lectures

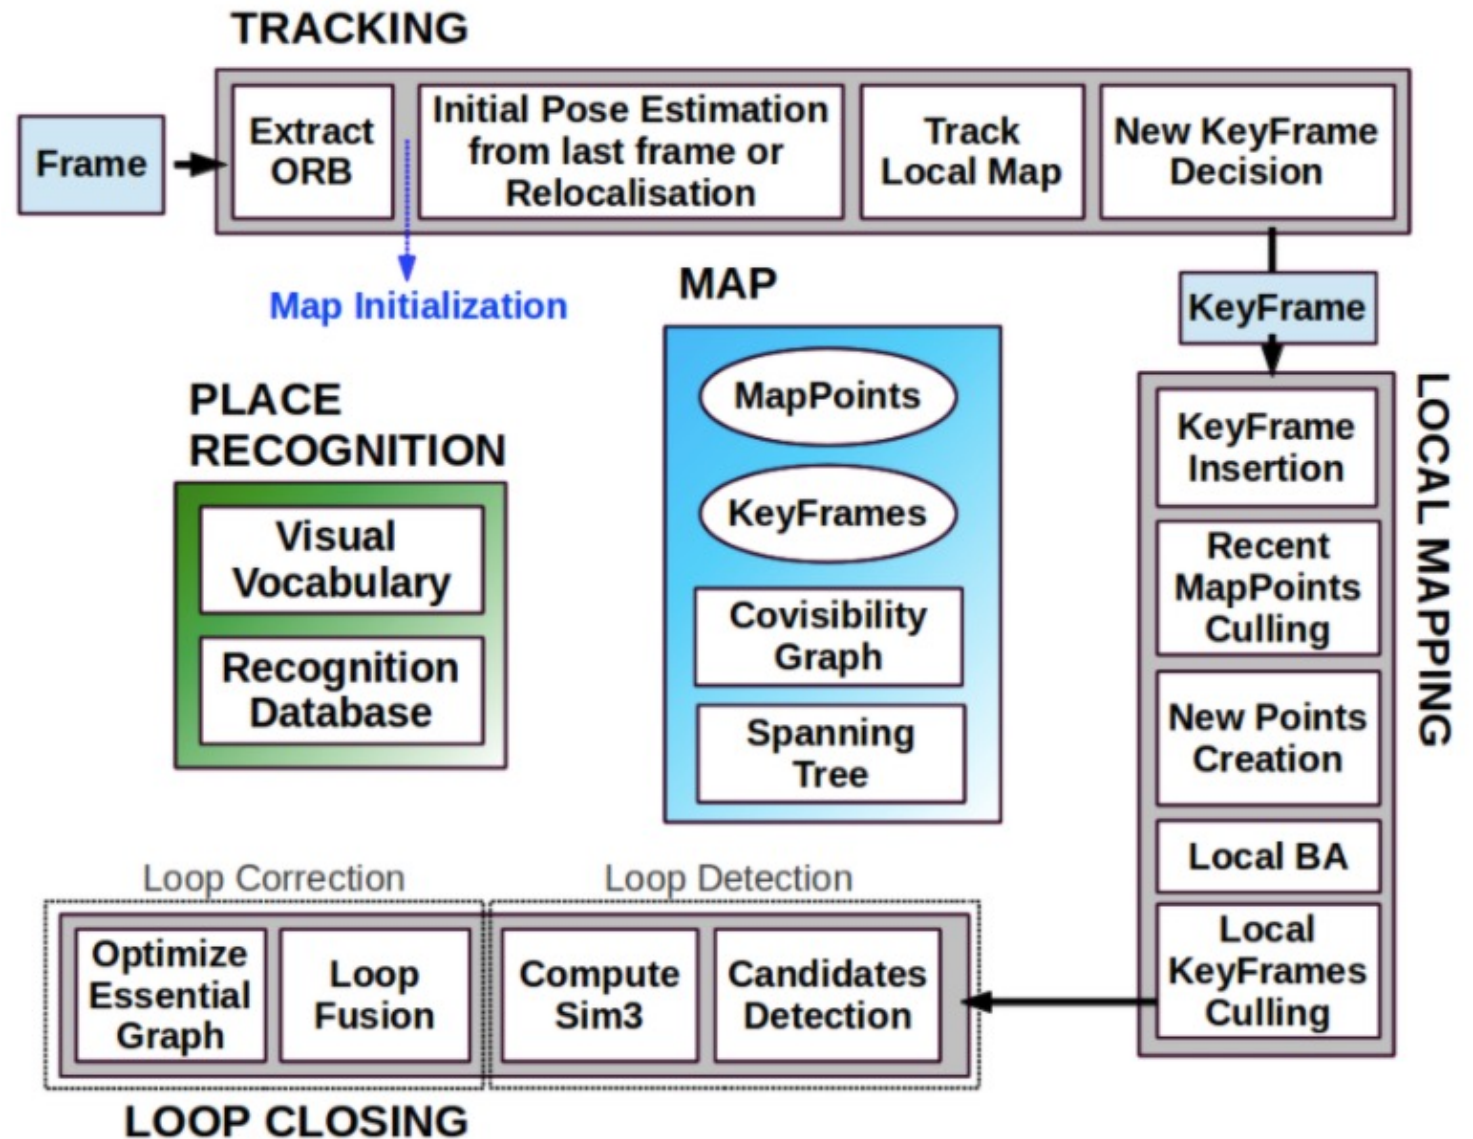


Feature-based Methods

- Feature Extraction
 - Feature Detectors & Feature Descriptor; more in vision lectures
 - ORB, SIFT, AKAZE, BRISK, etc ...
- Feature Matching
 - BFM, KNN, etc ...
- Relative Pose Calculation
 - 5-pt, 7-pt, 8-pt, PnP, etc ...


Feature-based Method: ORB-SLAM

Mur-Artal R, Montiel J M M, Tardos J D. ORB-SLAM: a versatile and accurate monocular SLAM system[J]. IEEE transactions on robotics, 2015, 31(5): 1147-1163.



Direct Method: LSD-SLAM

- Construct Photometric Error


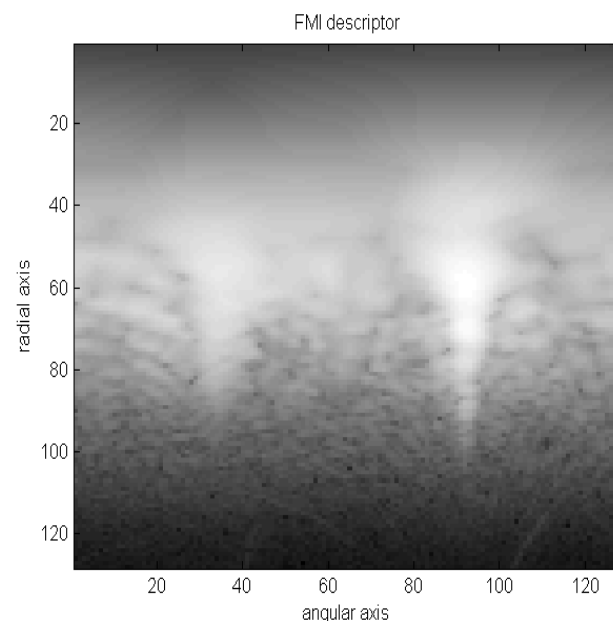



- Construct Depth Error
- Minimize Objective Error Function

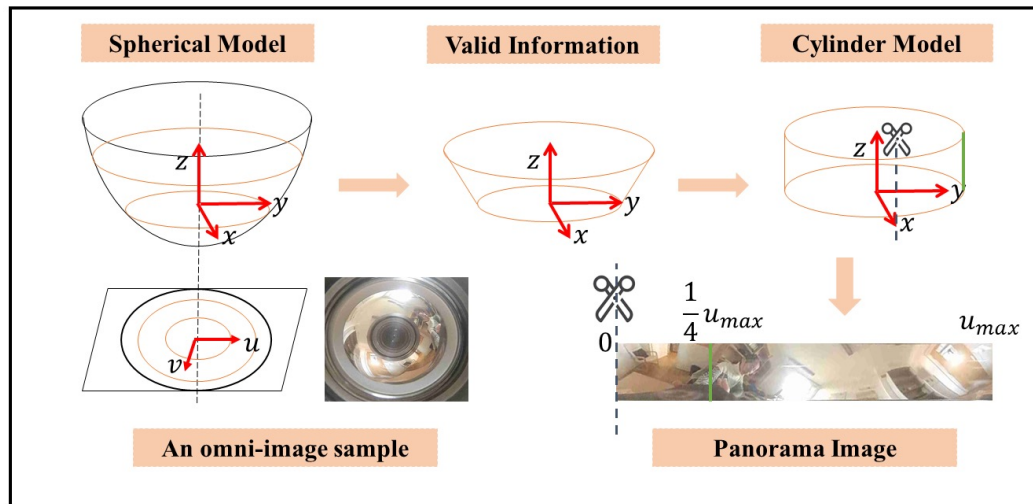
Engel J, Schöps T, Cremers D. LSD-SLAM: Large-scale direct monocular SLAM[C]//European conference on computer vision. Springer, Cham, 2014: 834-849.

Direct Method: Fourier Mellin Transform

- Spectral based registration: detection of scaling, rotation and translation in 2 subsequent frames
- Processing spectrum magnitude decouples translation from affine transformations
 - Detection of signal shift between 2 signals by phase information
 - Resampling to polar coordinates → Rotation turns into signal shift !
 - Resampling the radial axis from linear to logarithmic presentation → Scaling turns into signal shift !
 - Calculate a Phase Only Match Filter (POMF) on the resampled magnitude spectra




Pose Estimation for Omni-directional Cameras using Sinusoid Fitting



$$y = B + A \sin(\omega x + \phi)$$

$$\Delta v(u_p) = \lambda_p t_z + \gamma \left\| P_{xy}(R) \cdot \sin(\gamma u_p + \frac{P_{xy}(R)}{\|P_{xy}(R)\|}) \right\|$$


$$\Delta u(u_p) = \gamma P_z(R) + \lambda_p \|P_{xy}(t)\| \cdot \sin(\gamma u_p + \hat{t}_{xy})$$



(1a)

(1b)

(1c)



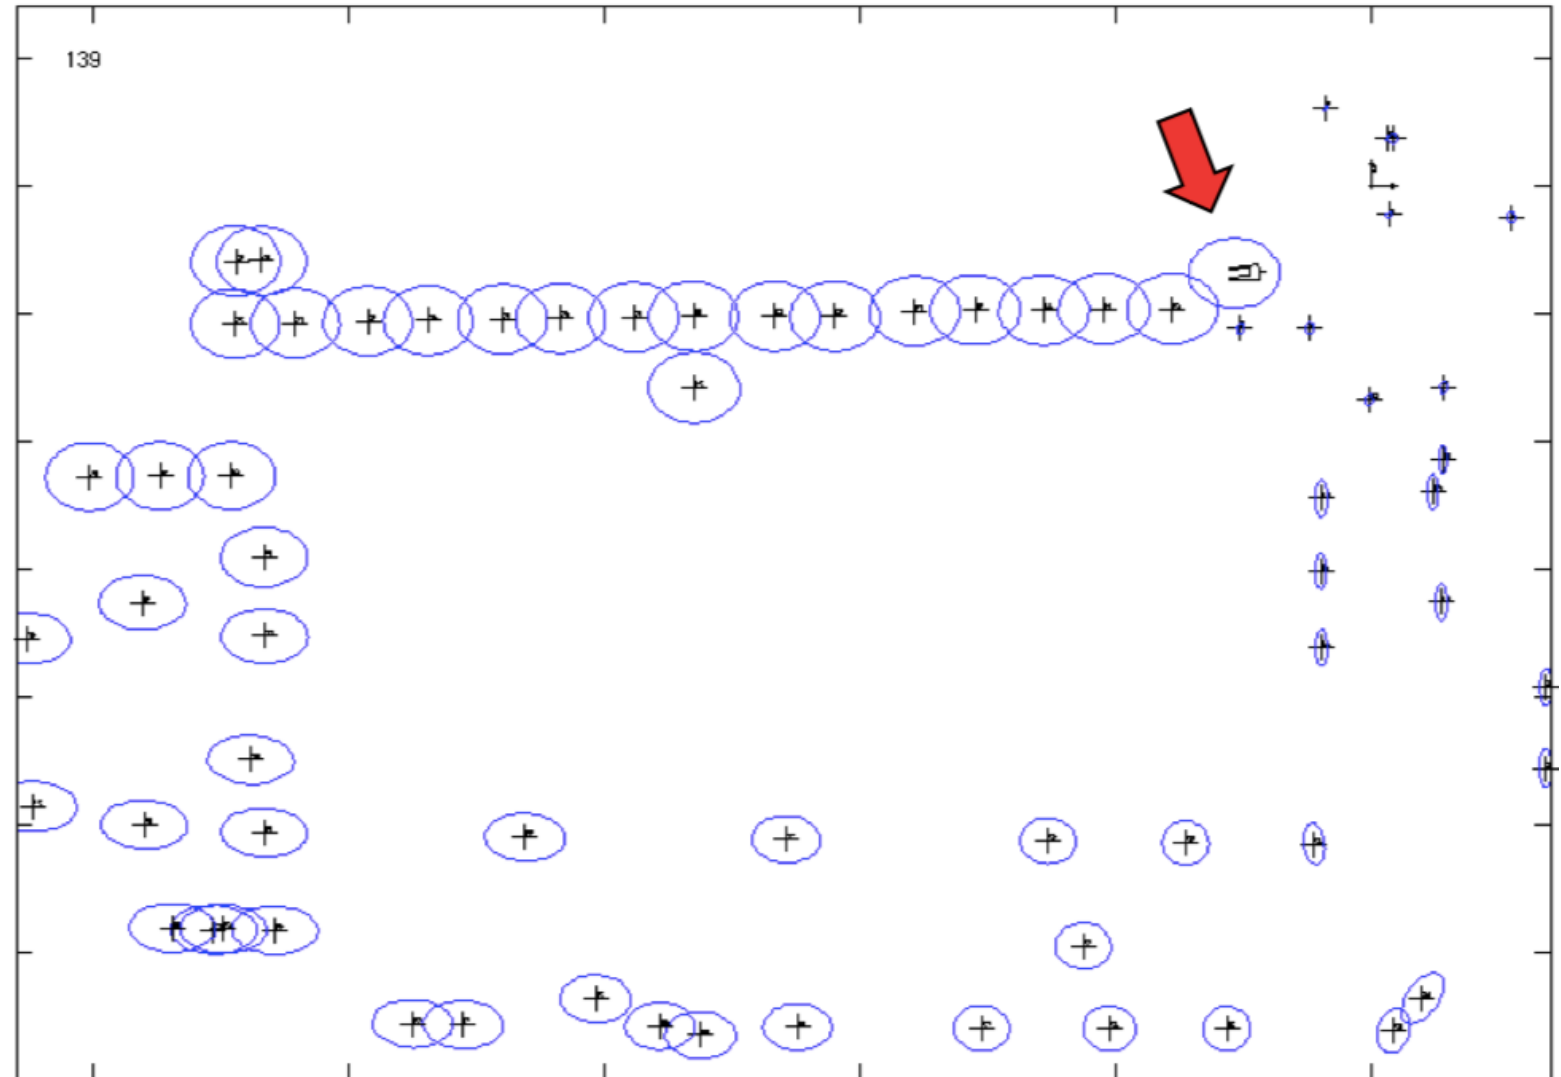
BACK END

Overview of Back-end

- Loop Detection
 - Find candidates of scan pairs/ scan with old map
 - E.g. based on global pose estimated (chain rule) OR image similarity (bag of words)
- Loop Closure
 - E.g. use scan matching to find the transform AND its uncertainty
- Optimization
 - Pose Graph optimization (e.g. minimize error of poses, based on uncertainty)
 - Bundle Adjustment
- Map Rendering
 - E.g. generate grid map based on optimized graph


Loop Closure

- Before loop closure



Loop Closure

- After loop closure



Loop Closure

- Recognizing an already mapped area, typically after a long exploration path (the robot “closes a loop”)
- Structurally identical to data association, but
 - high levels of ambiguity
 - possibly useless validation gates
 - environment symmetries
- Uncertainties collapse after a loop closure (whether the closure was correct or not)

Loop Closure

- By revisiting already mapped areas, uncertainties in robot and landmark estimates can be reduced
- This can be exploited when exploring an environment for the sake of better (e.g. more accurate) maps
- Exploration: the problem of where to acquire new information

Robust Loop Closing over Time for Pose Graph SLAM

Instituto de Investigación en
Ingeniería de Aragón (I3A)

Yasir Latif: ylatif@unizar.es
José Neira: jneira@unizar.es

Volgenau School of Engineering
George Mason University

César Cadena: ccadenal@gmu.edu

*This work was supported by Spanish DPI2009-13710 and DPI2009-07130,
by DGA-FSE (group T04), and by the US Army Research Office
(W911NF-1110476)*

OVERVIEW: THREE SLAM PARADIGMS

The Three SLAM Paradigms

- Most of the SLAM algorithms are based on the following three different approaches:
 - Extended Kalman Filter SLAM: (called EKF SLAM)
 - Particle Filter SLAM: (called FAST SLAM)
 - Graph-Based SLAM

EKF SLAM: overview

- **Extended state vector** y_t : robot pose x_t + position of all the features m_i in the map:

$$y_t = [x_t, m_0, \dots, m_{n-1}]^T$$

- Example: 2D line-landmarks, size of $y_t = 3+2n$: three variables to represent the robot pose + $2n$ variables for the n line-landmarks having vector components

$$(\alpha_i, r_i)$$

$$y_t = [x_t, y_t, \theta_t, \alpha_0, r_0, \dots, \alpha_{n-1}, r_{n-1}]^T$$

- As the robot moves and takes measurements, the state vector and covariance matrix are updated using the standard equations of the extended Kalman filter
- Drawback: EKF SLAM is computationally very expensive.